

STAT400 Notes

Table of contents

1. Basic Terms	1
2. Probability functions	2
2.1. Sigma algebra	2
2.2. Probability functions	2
3. Independent events	2
4. Random variables	2
5. Discrete distributions	2
5.1. Geometric distribution	2
6. Continuous distributions	2
6.1. Normal distribution	2
6.2. Gamma distribution	2
6.3. Exponential distribution	3
6.4. Chi-squared distribution	3
7. Relationships between variables	3
7.1. Covariance	3
7.2. Correlation coefficient	4
8. Linear combinations of random variables	4
9. Random samples and statistics	4
9.1. Central Limit Theorem	4
10. Estimators	5
10.1. Parametric estimation	5
10.1.1. Point estimators	5
10.2. Method of Moments	5
10.3. Maximum Likelihood Estimators (MLEs)	5

1. Basic Terms

Parameters Quantitative features of a population

Statistics Quantitative features calculated using a sample

Deduction A way of going from the population to the sample, e.g., if you get the class average (population) and then compare it to your personal score on an exam (sample), you can deduce that you did better than the average.

Inference A way of going from the sample to the population (opposite direction from deduction). For example, if you get a 100 on an exam (sample) and you infer that you did better than the class average (population).

Inference and deduction goes in a circle. When developing a model, you estimate some parameters, then you make deductions about the population, then you infer population parameters, then you draw deductions about the population again.

Experiment Repeatable task with well-defined outcomes

Sample space The sample space for an experiment is the set of all possible outcomes of that experiment, denoted by S

Event Any subset E of the sample space S attached to an experiment will be called an event associated with the experiment. To say that an event E has happened means that the outcome of the experiment was in the set E .

Simple event An event with only one event in it

2. Probability functions

2.1. Sigma algebra

A sigma algebra is a collection of interesting events in some sample space.

A collection \mathcal{B} of subsets of a sample space S is a sigma algebra if:

1. $\emptyset \in \mathcal{B}$
2. $\forall A \in \mathcal{B}, A^c \in \mathcal{B}$
3. If $\{A_i : i \in \mathbb{N}\}$ is a countable collection such that $A_i \in \mathcal{B}$ for all i , then $\bigcup_{i \in \mathbb{N}} A_i \in \mathcal{B}$

2.2. Probability functions

Consider a sample space S with a sigma algebra \mathcal{B} .

A probability function is a function from events to probabilities ($\mathcal{B} \rightarrow \mathbb{R}$). It must satisfy the following axioms:

1. (finite measure) $P(S) = 1$
2. (positivity) $\forall A \in \mathcal{B}, P(A) \geq 0$
3. (countable additivity) For A_1, A_2, A_3, \dots , the collection of pairwise disjoint subsets of S in \mathcal{B} , we must have

$$P\left(\bigcup_{i \in \mathbb{N}} A_i\right) = \sum_{i=1}^{\infty} P(A_i)$$

3. Independent events

Two events A and B are independent if any of the following are true (all are equivalent):

- $P(A \cap B) = P(A)P(B)$
- $P(A | B) = P(A)$
- $P(B | A) = P(B)$

4. Random variables

A random variable X maps outcomes in some sample space to real numbers, i.e., $X : \mathcal{S} \rightarrow \mathbb{R}$. A random variable measures a specific quantitative feature of the sample space outcome.

The **range** of X , the set of all possible values that X can take, is denoted \mathcal{X} .

With the new sample space, \mathcal{X} , you can use the order relationship in real numbers and you can add, multiply, etc.

5. Discrete distributions

5.1. Geometric distribution

6. Continuous distributions

6.1. Normal distribution

todo

6.2. Gamma distribution

The gamma distribution is based on the gamma function, which extends the factorial function to complex numbers:

$$\Gamma(\alpha) = \int_0^{\infty} x^{\alpha-1} e^{-x} dx$$

Some properties of Γ :

- $\Gamma(x) = x!$ if x is a non-negative integer
- $\Gamma(a+1) = a\Gamma(a)$

If you have a gamma distribution $\text{Gamma}(\alpha, \lambda)$, then its probability density function is

$$f_{X(x)} = \begin{cases} \frac{\lambda^\alpha x^{\alpha-1} e^{-\lambda x}}{\Gamma(\alpha)} & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

Note: there are two different parameterizations for Gamma:

- $\text{Gamma}(\alpha, \beta/\lambda)$ (with a rate parameter)
- $\text{Gamma}(\alpha, \theta)$ (with a scale parameter)

The α is the shape parameter. $\beta = \frac{1}{\theta}$

6.3. Exponential distribution

This is a special case of the gamma distribution: $\text{Exponential}(\lambda) = \text{Gamma}(1, \lambda)$. If you sum n independent $\text{Exponential}(\lambda)$ random variables, you get a $\text{Gamma}(n, \lambda)$ random variable.

Application: Variables that model the amount of time you have to wait before something happens follow an exponential distribution. E.g., time between clicks of a Geiger counter.

This is the continuous analog of the geometric distribution. Only the geometric and exponential distributions are **memoryless**.

6.4. Chi-squared distribution

Also a special case of the gamma distribution: $\chi_k^2 \sim \text{Gamma}(\alpha = \frac{k}{2}, \theta = 2)$.

Stuff in real life isn't distributed this way. Chi-squared is mostly just used for hypothesis tests. It's closely related to the standard normal distribution (Z): $\chi_1^2 = Z^2$. In general:

$$\chi_k^2 = (Z_1)^2 + (Z_2)^2 + \dots + (Z_k)^2$$

7. Relationships between variables

Large value A value of some random variable is large if it's greater than the mean

Small value A value of some random variable is small if it's less than the mean

Positive relationship X and Y have a positive relationship if large values of X are associated with large values of Y and small values of X are associated with small values of Y

Negative relationship X and Y have a negative relationship if large values of X are associated with small values of Y and small values of X are associated with large values of Y

7.1. Covariance

$$\text{Cov}(X, Y) = E((X - \mu_X)(Y - \mu_Y))$$

When X and Y have a positive relationship, the covariance should be positive. When they have a negative relationship, the covariance should be negative.

Units of $\text{Cov}(X, Y)$ are (units of X) \cdot (units of Y). This is a problem, since the magnitude of the covariance depends on the units.

Properties:

- $\text{Cov}(X, X) = V(X)$
- $\text{Cov}(X, Y) = \text{Cov}(Y, X)$
- $\text{Cov}(aX + bY, Z) = a\text{Cov}(X, Z) + b\text{Cov}(Y, Z)$

7.2. Correlation coefficient

$$\text{Corr}(X, Y) = \rho_{X,Y} = \frac{\text{Cov}(X, Y)}{\sqrt{V(X)}\sqrt{V(Y)}}$$

This one is unitless, unlike covariance. It always lies in $[-1, 1]$.

$\rho_{X,Y} = \pm 1$ iff there is a perfect linear relationship between X and Y .

The correlation coefficient measures the extent of the linear relationship between X and Y .

8. Linear combinations of random variables

If you have $Y = a_1X_1 + a_2X_2 + \dots + a_nX_n$, then $E(Y) = a_1E(X_1) + \dots + a_nE(X_n)$

Variance is more complicated:

$$V(Y) = \text{Cov}(Y, Y) = \sum_{i=1}^n a_i^2 V(X_i) + 2 \sum_{1 \leq i < j \leq n} a_i a_j \text{Cov}(X_i, X_j)$$

If all the random variables X_i are independent, then the covariance terms cancel out, meaning that:

$$V(Y) = \sum_{i=1}^n a_i^2 V(X_i)$$

9. Random samples and statistics

Random sample We say the collection of random variables $\{X_1, X_2, \dots, X_n\}$ is a **random sample** of size n from the population distribution X if:

- The X_i s are identically distributed to the distribution of X , i.e., $X_i \sim X$
- The X_i s are mutually independent, i.e., joint pdf/pmf splits into its marginals

Statistic A quantity calculated using a random sample.

Joint sample space Suppose $\{X_1, \dots, X_n\}$ is a random sample from population X and X takes values in \mathcal{X} . The joint sample space for the random sample is $\mathcal{X}^n = \{(x_1, \dots, x_n) \mid x_i \in \mathcal{X}\}$.

The joint sample space is the set of all possible sample data for the random sample.

If T is a statistic calculated using the random sample $\{X_1, \dots, X_n\}$ with joint sample space \mathcal{X}^n , then we can think of T as the function

$$T : \mathcal{X}^n \rightarrow \mathbb{R}$$

Therefore, T is a random variable.

Sampling distribution The associated probability distribution of T above is called the sampling distribution of the statistic T .

9.1. Central Limit Theorem

If you have a random sample of size n from a distribution X with $E(X) = \mu$ and $V(X) = \sigma^2$, then, if n is large enough, the sample mean \bar{X} will approximately have the distribution $N\left(\mu, \frac{\sigma^2}{n}\right)$

10. Estimators

10.1. Parametric estimation

Suppose the actual distribution has parameters $\theta_1, \theta_2, \dots, \theta_k$.

We want to estimate these parameters. The estimates are $\hat{\theta}_1, \hat{\theta}_2, \dots, \hat{\theta}_k$

10.1.1. Point estimators

Suppose a population is fixed and has distribution X with parameters $\theta_1, \theta_2, \dots, \theta_k$

Point estimator A point estimator for the parameter θ_i is a statistic $\hat{\theta}_i$ calculated using a random sample of size n coming from the population distribution whose values are used as estimates for θ_i

Bias The bias of a point estimator is the expected deviation of values from $\hat{\theta}$ from θ

$$\text{Bias}(\hat{\theta}) = E(\hat{\theta}) - \theta = E(\hat{\theta} - \theta)$$

Unbiased A point estimator $\hat{\theta}$ is unbiased if $\text{Bias}(\hat{\theta}) = 0$ for all θ

Consistent A point estimator $\hat{\theta}$ is consistent if $\lim_{n \rightarrow \infty} \text{MSE}(\hat{\theta}) = 0$

Very important theorem:

$$\text{MSE}(\hat{\theta}) = V(\hat{\theta}) + \text{Bias}(\hat{\theta})^2$$

10.2. Method of Moments

Goal: Given k parameters attached to a population distribution. and a random sample of size n , find k estimators for these parameters.

Intuition: the estimators will show up as solutions to a system of k equations.

Population moment The k th population moment is defined as $\mu_k := E(X^k)$

Sample moment Given a random sample $\{X_1, X_2, \dots, X_n\}$ coming from the population distribution X , the k th sample moment is defined as $\frac{1}{n} \sum_{i=1}^n X_i^k$

Intuition: For large enough n , the k th sample moments estimate the k th population moments.

MoM Algorithm:

1. Calculate the first k population moments
2. Derive a system of k equations by equating each i th population moment to the i th sample moment.
3. Find solutions to the system, if they exist. These are the MoM estimators.

10.3. Maximum Likelihood Estimators (MLEs)

Goal is same as Method of Moments estimators: given k parameters attached to a population distribution. and a random sample of size n , find k estimators for these parameters.

Intuition: set up an optimization problem, and the solutions to this problem, if they exist, are the MLEs.

Suppose the population has distribution X with pmf given by $f(x; \theta_1, \theta_2, \dots, \theta_k)$

Suppose $\{X_1, X_2, \dots, X_n\}$ is sample data coming from X

Likelihood function (you want to maximize this):

$$L : \Theta \rightarrow \mathbb{R}$$

$$L(\theta_1, \dots, \theta_k; X_1, \dots, X_n) := \prod_{i=1}^n f(X_i; \theta_1, \dots, \theta_k)$$

($\Theta \subseteq \mathbb{R}^k$ is the parameter space)

Log likelihood function: $l(\theta_1, \dots, \theta_k; X_1, \dots, X_n) = \ln(L(\theta_1, \dots, \theta_k; X_1, \dots, X_n))$

The arguments that maximize the log likelihood function also maximize the likelihood function, so you can try to find the maximize the log likelihood function instead of the likelihood function. This is helpful because maximizing the log likelihood function is easier.

Principal of maximum likelihood estimation: If $\hat{\theta}_1, \dots, \hat{\theta}_k$ satisfy $L(\hat{\theta}_1, \dots, \hat{\theta}_k; X_1, \dots, X_n) \geq L(\theta_1, \dots, \theta_k; X_1, \dots, X_n)$ for all $(\theta_1, \dots, \theta_k) \in \mathbb{R}^k$, then $\hat{\theta}_1, \dots, \hat{\theta}_k$ are called the **maximum likelihood estimators** for $\theta_1, \dots, \theta_k$.

To calculate MLE, apply second derivative test to likelihood function.